

Accelerate JAX on Intel GPUs via PJRT

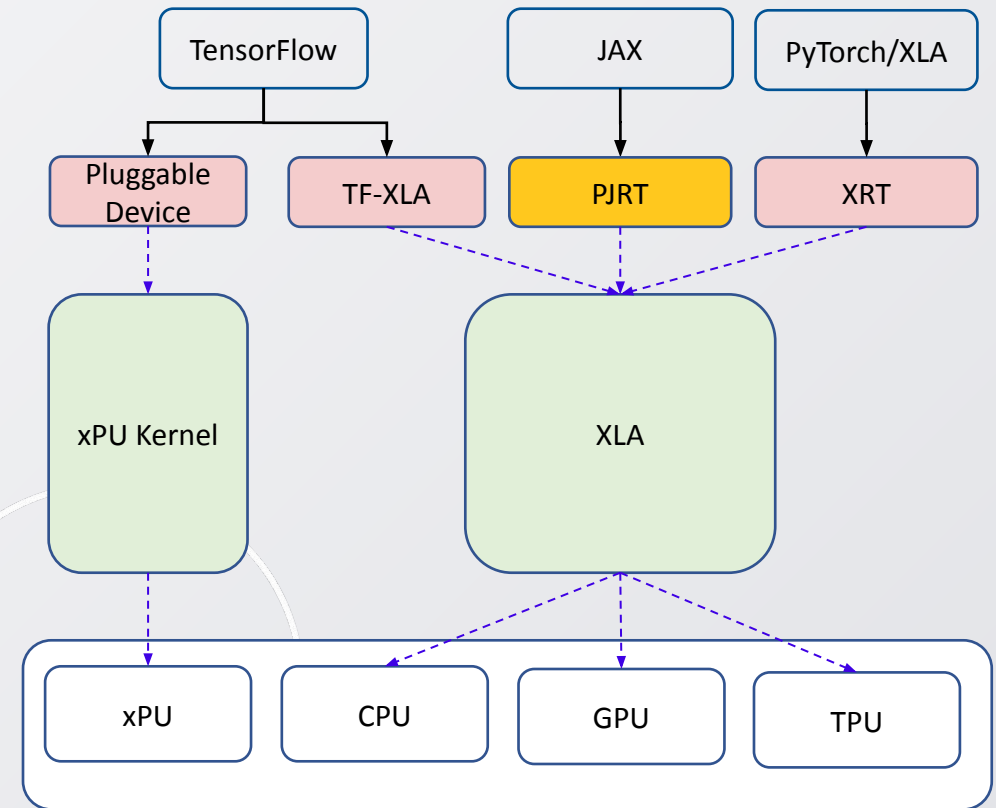
Xiao Yu, Google

Yiqiang Li, Intel



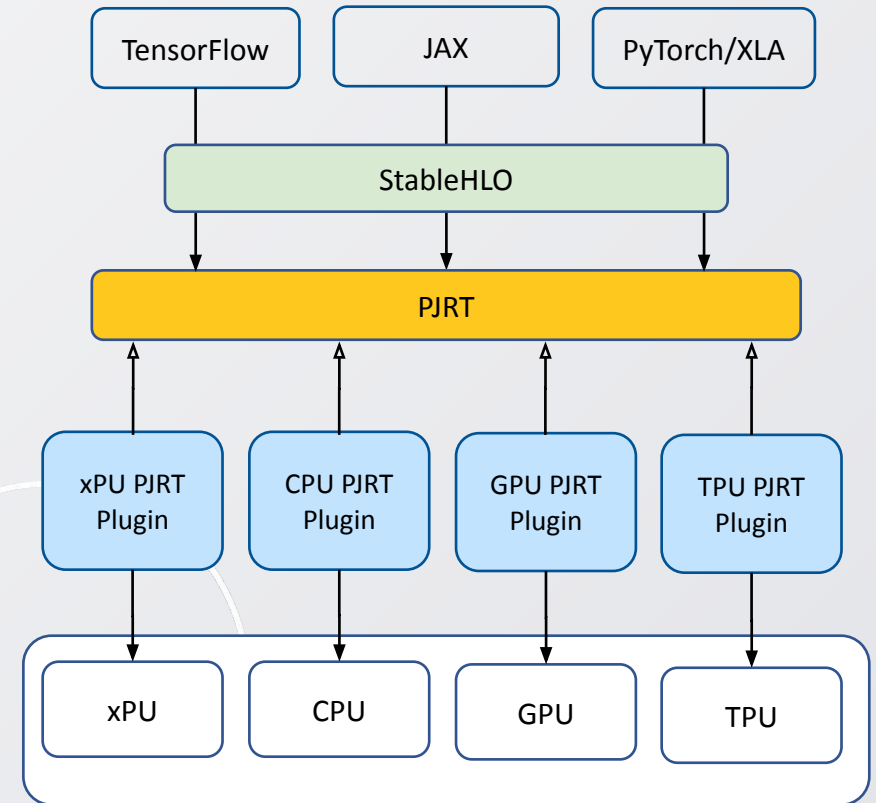
Problem

- Fragmented stack for different ML frameworks
- No way to add new hardware support in JAX
- TensorFlow has its own way for adding new hardware support
 - Cannot leverage compiler technologies



PJRT

- Unified pluggable device API --- **PJRT**
- A PJRT Plugin contains:
 - Hardware-specific compiler, which takes StableHLO as standard input IR.
 - Hardware runtime
- Decouple ML Framework and device PJRT plugin
 - Plugin are discovered and loaded as dynamic library

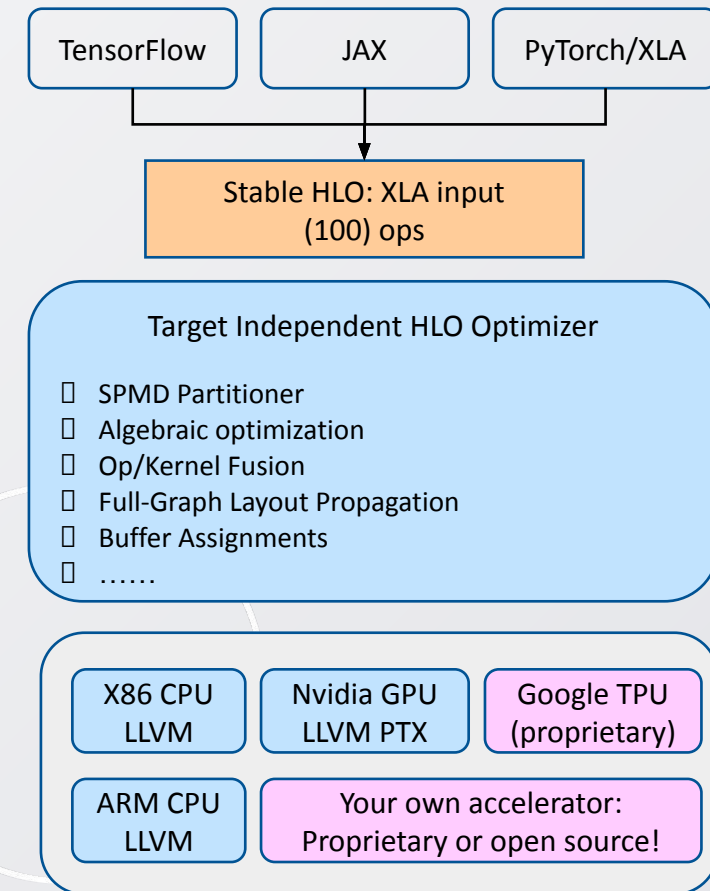


PJRT: high level overview

- Scope: “**Compile and Execute a StableHLO program**”
- A PJRT Plugin will implement following APIs:
 - **Compile:** StableHLO -> PjrtLoadedExecutable
 - Trigger HW specific compiler to compile StableHLO into a HW executable
 - **H2D transfer:** host buffer, PjrtDevice -> PjrtBuffer
 - Prepare input by transferring data from host to HW
 - **Execute:** PjrtLoadedExecutable, PjrtDevice, PjrtBuffer -> PjrtBuffer
 - Execute the program on HW
 - **D2H transfer:** PjrtBuffer -> host buffer
 - Transfer output back to host

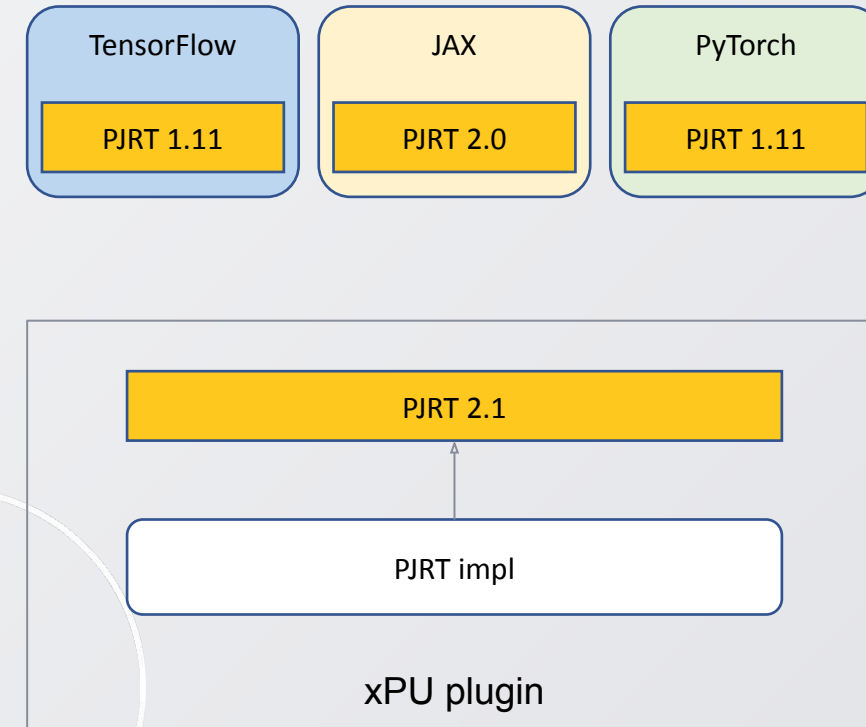
OpenXLA

- It is challenging to build a new device compiler and runtime from scratch
- OpenXLA: Open, state-of-the-art ML compiler, using the best of XLA & MLIR
- Reusable target independent optimizer




Versioning

- ML Frameworks and PJRT Plugins can be released separately
 - ML Frameworks and PJRT Plugins may use different version of PJRT.
- Proposal:
 - The program works as long as the framework and the plugin have the same Major version.



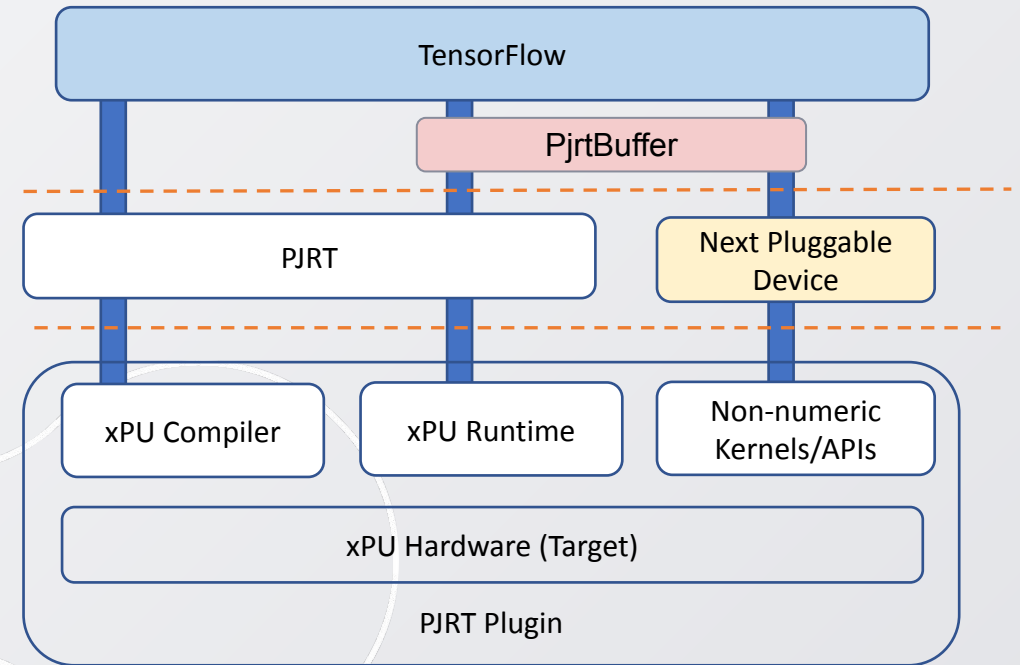
Ahead of Time Compilation

- Compile and execute StableHLO program on different machines
- Supported by Jax (done) and TensorFlow (WIP)
- PJRT Plugin is responsible for defining versioning and compatibility of serialized executable



Custom Kernel support

- Custom Kernels are critical to support
 - Operation that cannot be represented in StableHLO.
 - Hand-written kernel
- Use PjrtBuffer as the concurrency to enable efficient (0 Copy) buffer sharing between PJRT and custom kernels
 - Implemented in TensorFlow as NextPluggableDevice API

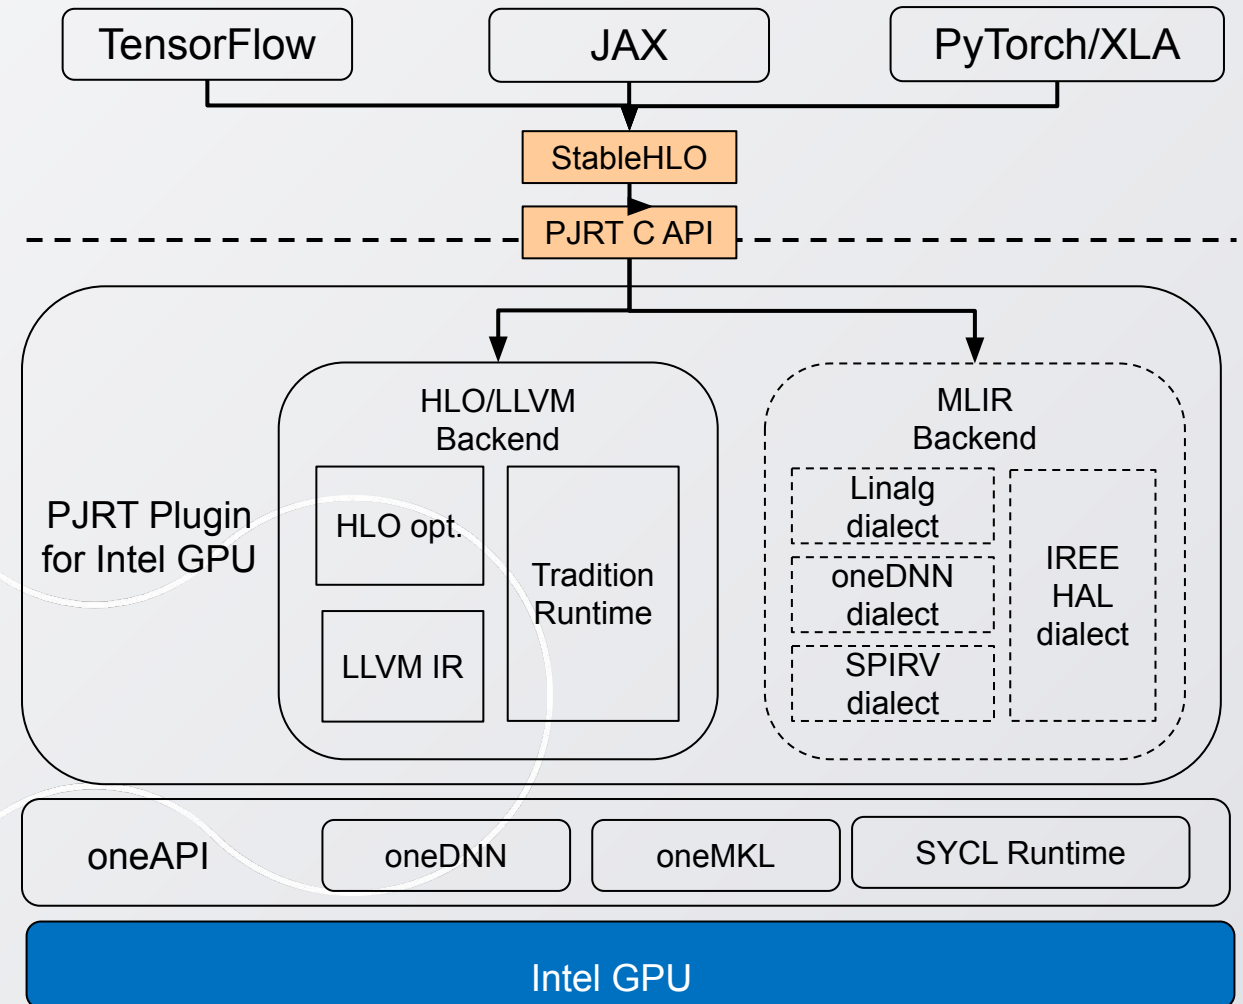


PJRT - Future

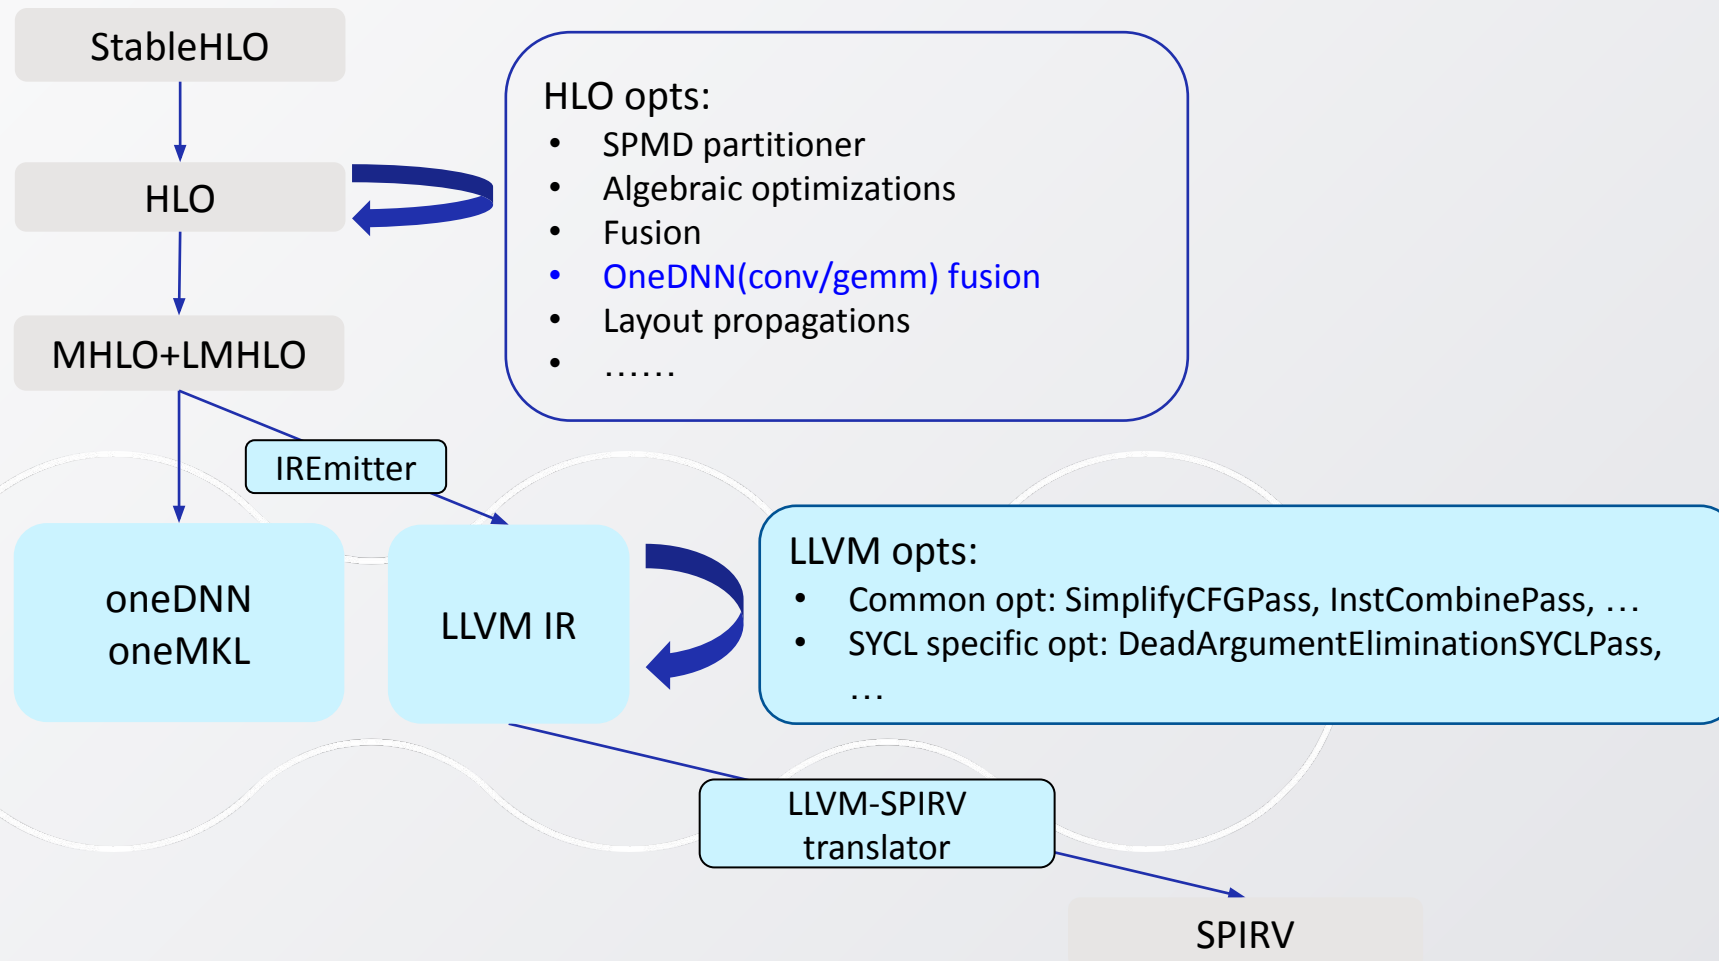
- Memory Space Support
 - Example use case: Nvidia Grace Hopper Superchip
- Sparsity Support
 - Enable ML compiler to optimize sparse computation
- MPMD Support
 - Enable more advanced Pipelining Parallelism

PJRT Plugin for Intel GPU

- Intel GPU plugin integrates via PJRT API using oneAPI, and works on two code-gen techniques in parallel:
 - HLO/LLVM IR, release, runs JAX models
 - MLIR/IREE, experimenting, runs simple cases
- Release to support JAX ([link](#))



PJRT Plugin for Intel GPU



PJRT Plugin for Intel GPU --- Library Opt

- oneDNN
 - Conv/GEMM
 - Supported fusion pattern
 - Conv + [bias + add + activation]
 - [alpha *] GEMM (A, B) + [beta * C]
 - Graph API integration is coming soon to have more fusion capabilities
- oneMKL: Cholesky, FFT, TriangularSolve

Python code:

```
@jax.jit
def lax_conv_example(lhs, rhs):
    out = jax.lax.conv_with_general_padding(
        lhs, rhs, (1,1), ((0,0),(0,0)), (1,1), (1,1))
    out = jax.nn.relu(out)
    return out

lhs = lhs = np.random.randn(2,1,9,9).astype(np.float32)
rhs = np.random.randn(1,1,4,4).astype(np.float32)
lax_conv_example(lhs, rhs)
```

StableHLO:

```
module @jit_lax_conv_example {
  func.func public @main(%arg0: tensor<2x1x9x9xf32>, %arg1: tensor<1x1x4x4xf32>) ->
  (tensor<2x1x6x6xf32>) {
    %0 = stablehlo.convolution(%arg0, %arg1) : (tensor<2x1x9x9xf32>,
    tensor<1x1x4x4xf32>) -> tensor<2x1x6x6xf32>
    %1 = call @relu(%0) : (tensor<2x1x6x6xf32>) -> tensor<2x1x6x6xf32>
    return %1 : tensor<2x1x6x6xf32>
  }
  func.func private @relu(%arg0: tensor<2x1x6x6xf32>) -> tensor<2x1x6x6xf32> {
    %0 = stablehlo.constant dense<0.000000e+00> : tensor<f32>
    %1 = stablehlo.broadcast_in_dim %0, dims = [] : (tensor<f32>) ->
    tensor<2x1x6x6xf32>
    %2 = stablehlo.maximum %arg0, %1 : tensor<2x1x6x6xf32>
    return %2 : tensor<2x1x6x6xf32>
  }
}
```

Optimized HLO:

```
ENTRY main.10 {
  constant_3 = f32[1]{0} constant({0})
  Arg_1.2 = f32[1,1,4,4]{3,2,1,0} parameter(1)
  Arg_0.1 = f32[2,1,9,9]{3,2,1,0} parameter(0)
  onednn-conv-bias-activation.1 = (f32[2,1,6,6]{3,2,1,0}, u8[0]{0})
  custom-call(Arg_0.1, Arg_1.2, constant_3), backend_config="{\"activation_mode\": \"2\"}"
  -> conv+relu fusion
  ROOT get-tuple-element = f32[2,1,6,6]{3,2,1,0}
  get-tuple-element(onednn-conv-bias-activation.1), index=0
}
```

PJRT Plugin for Intel GPU --- LLVM IR

Python code:

```
@jax.jit
def func_jit(a):
    a =
    jnp.abs(jnp.sqrt(x))
    return a
```



StableHLO module:

```
module @jit_func {
  func.func public @main(%arg0:
    tensor<1024xf32> {jax.arg_info = "x", ...}) {
    %0 = stablehlo.sqrt %arg0 :
    tensor<1024xf32> loc(#loc2)
    %1 = stablehlo.abs %0 : tensor<1024xf32>
    loc(#loc3)
    return %1 : tensor<1024xf32> loc(#loc)
  } loc(#loc)
} loc(#loc)
```



HLO module:

```
%fused_computation (param_0.1: f32[1024]) -> f32[1024] {
  %param_0.1 = f32[1024]{0} parameter(0)
  %sqrt.0 = f32[1024]{0} sqrt(f32[1024]{0} %param_0.1),
  metadata={op_type="Sqrt" ...}
  ROOT %abs.0 = f32[1024]{0} abs(f32[1024]{0} %sqrt.0),
  metadata={op_type="Abs" ...}
}
ENTRY %func.8 (arg0.1: f32[1024]) -> f32[1024] {
  %arg0.1 = f32[1024]{0} parameter(0), ...
  ROOT %fusion = f32[1024]{0} fusion(f32[1024]{0} %arg0.1),
  kind=kLoop, calls=%fused_computation,
}
```



Lmhlo module:

```
module attributes {hlo.unique_id = 0 : i32, mhlo.unique_id = 0 : i64} {
  func @func(%arg0: memref<4096xi8> {lmhlo.params = 0 : index}, %arg1:
    memref<4096xi8> {lmhlo.output_index = dense<> : tensor<0xi64>})
  attributes {result_xla_shape = "f32[1024]{0}" } {
    %c0 = arith.constant 0 : index
    %0 = memref.view %arg0[%c0][] : memref<4096xi8> to memref<1024xf32>
    %c0_0 = arith.constant 0 : index
    %1 = memref.view %arg1[%c0_0][] : memref<4096xi8> to
    memref<1024xf32>
    "lmhlo.fusion"() ({
      %2 = bufferization.to_tensor %0 : memref<1024xf32>
      %3 = "mhlo.sqrt"(%2) : (tensor<1024xf32>) -> tensor<1024xf32>
      %4 = "mhlo.abs"(%3) : (tensor<1024xf32>) -> tensor<1024xf32>
      memref.tensor_store %4, %1 : memref<1024xf32>
      "lmhlo.terminator"() : () -> ()
    }) : () -> ()
    "lmhlo.terminator"() : () -> ()
  }
}
```

PJRT Plugin for Intel GPU --- LLVM IR

Difference with NVVM:

- Target data layout and triple
- Address space
- SPIRV builtin function
 - get_global_id, ...
 - Subgroup shuffle, barrier
 - Math function: sqrt, expm

Address Space	NVPTX Memory Space	SPIR-V Memory Space
0	Generic	Private
1	Global	Global
2	*Internal Use	Constant
3	Shared	Workgroup
4	Constant	Generic
5	Local	

LLVM IR:

```
target datalayout = "e-p:64:64:64-i1:8:8-i8:8:8-xxxxx"
target triple = "spir64-unknown-unknown"
```

```
define spir_func void @fusion(i8 addrspacel1)* noalias nocapture
readonly align 16 dereferenceable(4096) %alloc0, i8 addrspacel1)*
noalias nocapture writeonly align 128 dereferenceable(4096)
%alloc1) local_unnamed_addr !intel_reqd_sub_group_size !1 {
entry:
  %0 = call i64 @_Z12get_group_idj(i32 0)
  %block_id = trunc i64 %0 to i32
  %1 = call i64 @_Z12get_local_idj(i32 0)
  %thread_id_x = trunc i64 %1 to i32
  %2 = shl nuw nsw i32 %block_id, 10
  %linear_index = add nuw nsw i32 %2, %thread_id_x
  %3 = bitcast i8 addrspacel1)* %alloc0 to float addrspacel1)*
  %4 = zext i32 %linear_index to i64
  %5 = getelementptr inbounds float, float addrspacel1)* %3, i64 %4
  %6 = load float, float addrspacel1)* %5, align 4, !invariant.load
!2
  %7 = call float @_Z4sqrtf(float %6)
  %8 = call float @llvm.fabs.f32(float %7)
  %9 = bitcast i8 addrspacel1)* %alloc1 to float addrspacel1)*
  %10 = getelementptr inbounds float, float addrspacel1)* %9, i64
%4
  store float %8, float addrspacel1)* %10, align 4
  ret void
}
declare spir_func float @_Z4sqrtf(float) local_unnamed_addr #0
```


Flax/JAX Stable Diffusion

<https://huggingface.co/CompVis/stable-diffusion-v1-4#jaxflax>

No code change is required for Intel GPU

```
import jax, sys, time
import numpy as np
from flax.jax_utils import replicate
from flax.training.common_utils import shard
from diffusers import FlaxStableDiffusionPipeline

scheduler, scheduler_state =
FlaxDPMSolverMultistepScheduler.from_pretrained("CompVis/stable-diffusion-v1-4",
subfolder="scheduler")
pipeline, params = FlaxStableDiffusionPipeline.from_pretrained("CompVis/stable-diffusion-v1-4",
scheduler=scheduler, revision="bf16", dtype=jax.numpy.bfloat16)
params["scheduler"] = scheduler_state
prompt = "a photo of an astronaut riding a horse on mars"

prng_seed = jax.random.PRNGKey(0)
prompt = jax.device_count() * [prompt]
prompt_ids = pipeline.prepare_inputs(prompt)

params = replicate(params)
prng_seed = jax.random.split(prng_seed, jax.device_count())
prompt_ids = shard(prompt_ids)

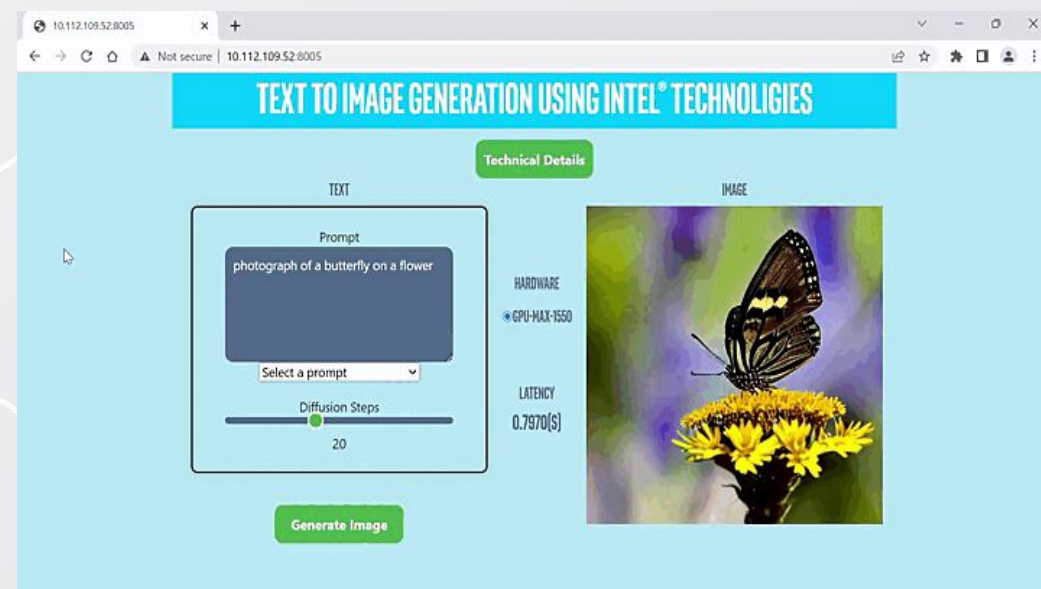
def elapsed_time(nb_pass=10, num_inference_steps=20):
    # warmup
    images = pipeline(prompt_ids, params, prng_seed, num_inference_steps, jit=True).images
    start = time.time()
    for _ in range(nb_pass):
        _ = pipeline(prompt_ids, params, prng_seed, num_inference_steps, jit=True).images
    end = time.time()
    return (end - start) / nb_pass

print("Latency per image is: {:.3f}s".format(elapsed_time(nb_pass=5, num_inference_steps=20)))
```

```
$ export LD_LIBRARY_PATH="python-path/jaxlib:$LD_LIBRARY_PATH"
$ export PJRT_NAMES_AND_LIBRARY_PATHS="xpu:/path/libitex_xla_extension.so"
$ numactl -N 0 -m 0 python jax_stable.py
```

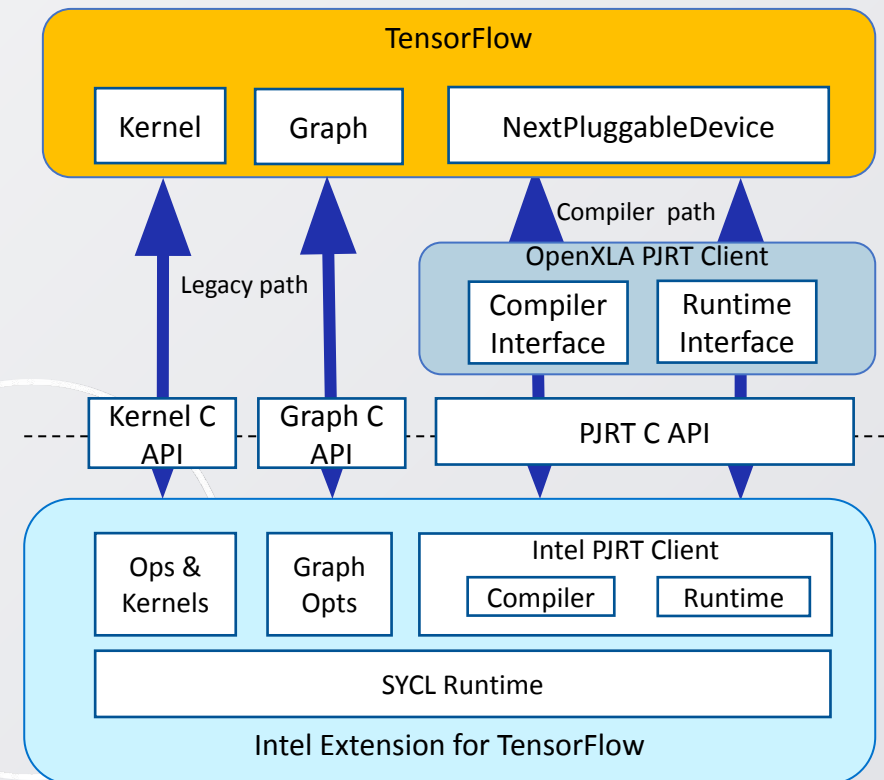
Latency per image is: **0.79s** on Intel® Data Center Max GPU 1550

Model	Precision	Diffusion Steps	Latency (s)
stable-diffusion-v 1-4	BF16	20	0.79
	BF16	50	1.84



NextPluggableDevice

- TensorFlow supports mixed execution mode
 - Run single TF model with both traditional and OpenXLA runtime
- Problems for TensorFlow plugins
 - PluggableDevice is StreamExecutor based, while OpenXLA is PJRT based
 - “NO” interoperability between them
- NextPluggableDevice solves this issue by extending PluggableDevice to use unified PJRT runtime for both



Summary

- PJRT simplifies ML Hardware and Framework integration with unified API to support all frameworks (TensorFlow, JAX, PyTorch via PyTorch-XLA)
- Intel GPU plugin integrated with JAX via PJRT API using oneAPI and LLVM/ SPIR-V and demonstrates good performance on Intel® Data Center Max GPU

Questions