


oneIPL Technical Advisory Board

Tech session #1

December 16th, 2021



oneAPI

Agenda

- Introduction – all TAB members (10 min)
- Technical discussion (45 min)
 - oneIPL overview:
 - Programming model
 - Execution model
 - Image processing pipelines
 - Image data abstraction
 - Memory model
- Closing words and next plans (5 min)

<https://spec.oneapi.io/oneipl/latest/index.html> - oneIPL specification (current version: v0.5)

The onePL TAB rules


DO NOT share any confidential information or trade secrets with the group

DO keep the discussion at a High Level

- Focus on the specific Agenda topics
- We are asking for feedback on features for the onePL specification (e.g. requirements for functionality and performance)
- We are NOT asking for the feedback on any implementation details

Please submit the feedback in writing on GitHub in accordance to [Contribution Guidelines](https://spec.oneapi.io/contributing) at spec.oneapi.io. This will allow Intel to further upstream your feedback to other standards bodies, including The Khronos Group SYCL specification.

Introduction of TAB members



oneAPI

- **Robert Schneider (PhD),**
Principal Key Expert
Diagnostic Imaging
Siemens Healthiness
- **SungShik Baik,**
Principle Engineer, PC engineer
Ultrasound System R&D
Samsung Medison
- **Kangsik Kim,**
Principle Engineer,
Ultrasound signal processing architect
Ultrasound System R&D
Samsung Medison
- **Ashish Uthama,**
Principal Software Engineer
Image Processing
Mathworks
- **Mark Rabotnikov,**
Lead software engineer,
Advanced Development group,
Enterprise Diagnostics Informatics
Philips
- **Tim van der Horst,**
C++ Software Designer,
Interventional Guided Therapy
Systems R&D - Imaging & Image
Processing
Philips
- **Sohrab Amirghodsi,**
Principal Compute Scientist
Photoshop ART
Adobe
- **Guoyi Zhou,**
Head of the Medical Innovation
Research Center
SonoScape
- **Zhilei Zhu,**
Xinje
- **Victor Getmanskiy,**
oneAPI architect,
Intel Performance Libraries,
Intel
- **Maksim Shabunin,**
AI Framework Engineer,
OpenVINO Core Engineering / OpenCV,
Intel
- **Sergey Ivanov,**
AI Framework Engineer,
OpenCV/G-API,
Intel

oneAPI Image Processing Library (oneIPL)

Domains overview



Currently available in the specification

Can be considered in the future versions of the specification

Basics

Image

Allocators

Accessors

Type conversions

Batch processing

Color conversions

RGB(A) \leftrightarrow NV12

RGB(A) \leftrightarrow RGBP

RGBP \leftrightarrow NV12

RGBA \leftrightarrow RGB

Other formats

Filters

Sobel 3 \times 3

Gaussian

Bilateral

Median

Other filters

Geometry

Resize

Mirror

Warp affine

Warp perspective

Other transforms

3D operations

Resize 3D

Remap 3D


Warp affine 3D

Median filter 3D

Other filters 3D

onePL programming language

- [SYCL 2020](#) – based on [C++17](#)
- onePL primitives - class data abstractions + functional API
- API shall be compatible with [SYCL 2020](#) compliant compiler implementation




oneAPI

onePL programming model

- Image data, layout, [region of interest \(ROI\)](#) are specified in [iopl::image class](#). Layout, data type, and memory are defined at compile-time.
- The supported [image formats](#) and data types are defined by [the matrix of combinations](#), each algorithm in the specification contain such matrix.
- Generic layouts is channel count – rows (1,3,4 channels). They are mapped to the formats: 1 – plane or grayscale, 3 – RGB, BGR, 4 – RGBA, BGRA, ...
- Additional layouts supported selectively – 3 planes for R, G, B, subsampled YUV formats (like NV12), etc.
- Generic datatypes – 8u-32u – unsigned integer, 8s-32s – signed integer, fp16-fp64 – floating-point


Layout	8u	8s	16u	16s	32u	32s	64u	64s	fp16	fp32	fp64	
1 channel												33 generic image formats
3 channels												
4 channels												
3 planes												extra formats
NV12		N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	



onePL programming model

- Each function has own type and format combination support matrix in the specification.
- Generic formats are usually supported for multiple devices.
- Some data types are device specific (FP16, FP64). Specification doesn't have the requirement that such API shall work as fallback. For example, if FP16 is not supported on device, it shall not be implemented via FP32. For that case error reporting is used.
- Example: resize_lanczos (support multichannel, NV12, doesn't support 3 planes)

Layout	8u	8s	16u	16s	32u	32s	64u	64s	fp16	fp32	fp64
1 channel	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
3 channels	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
4 channels	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
3 planes	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗
NV12	✓	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A




oneAPI

onePL programming model

- Image data, layout, [region of interest \(ROI\)](#) are specified in [ipl::image class](#). Layout, data type, and memory are defined at compile-time.
- Algorithm is a function called `<function>` which takes input image and output image. Additional arguments are passed as runtime class called `<function>_spec`. API has fixed arguments; spec structure can be extended.
- Example of [Gaussian algorithm API](#)

```
template <typename ComputeT = float,  
          formats Format,  
          typename DataT,  
          typename SrcAllocatorT,  
          typename DstAllocatorT>  
sycl::event gaussian(sycl::queue& queue,  
                    image<Format, DataT, SrcAllocatorT>& src,  
                    image<Format, DataT, DstAllocatorT>& dst,  
                    const gaussian_spec_t<Format, DataT, ComputeT>& spec,  
                    const std::vector<sycl::event>& dependencies = {})
```




oneAPI

onePL programming model

- ComputeT – type in which computations are done. E.g., implementation can provide faster solution with less precision using [sycl::half](#) rather than double data type.
- Each function takes [sycl::queue](#) and vector of dependencies.
- Function returns [sycl::event](#) which can be used in dependencies vector of other calls.

```
template <typename ComputeT = float,  
          formats Format,  
          typename DataT,  
          typename SrcAllocatorT,  
          typename DstAllocatorT>  
sycl::event gaussian(sycl::queue& queue,  
                    image<Format, DataT, SrcAllocatorT>& src,  
                    image<Format, DataT, DstAllocatorT>& dst,  
                    const gaussian_spec_t<Format, DataT, ComputeT>& spec,  
                    const std::vector<sycl::event>& dependencies = {})
```



oneAPI


onePL execution model

- onePL APIs follows [SYCL](#) xPU ideology. Each API is able to execute on the range of the devices, if device is supporting required features.
- onePL provides primitive algorithms so there is no covering of unsupported paths by different implementation. onePL algorithm shall not substitute unsupported type by the different supported type if it impacts the result. If the result precision doesn't matter, user still can write universal API covering all paths depending on device, so the ideology is that API shall be minimal for the range of the devices.
- Example: 2 devices CPU and GPU – float (fp32) is supported, double (fp64) is not supported on GPU, [sycl::half](#) (fp16) is supported on GPU.

```
(void)resize\_lanczos<float>(queue, input_image, output_image);
```

```
(void)resize\_lanczos<sycl::half>(queue, input_image, output_image); // valid only for GPU
```

```
(void)resize\_lanczos<double>(queue, input_image, output_image); // valid only for CPU
```




onePL execution model

- Execution mode is asynchronous. Algorithms submitted to [sycl::queue](#) and returns the control flow. Execution is scheduled by runtime taking into account the dependencies vector.
- For arguments having type [ipl::image](#) dependencies are handled automatically. In example below execution is sequential [convert](#) → [sobel 3x3](#) → [resize lanczos](#), since output of previous function is an input of next function.

```
// Source RGBA image data and destination images
image_t<formats::rgba, shared_usm_allocator_t> src_image{ queue,
                                                         src_image_data.get_pointer(),
                                                         src_size,
                                                         shared_allocator };

// Gray image data
image_t<formats::plane, device_usm_allocator_t> gray_image{ src_size, device_allocator };
// Sobel image data
image_t<formats::plane, device_usm_allocator_t> sobel_image{ src_size, device_allocator };
// Resized image data
image_t<formats::plane, shared_usm_allocator_t> resized_image{ dst_size, shared_allocator };

// Run pipeline: convert to grayscale -> sobel -> resize_lanczos
(void)convert(queue, src_image, gray_image);
(void)sobel_3x3(queue, gray_image, sobel_image);
(void)resize_lanczos(queue, sobel_image, resized_image);
```



oneAPI

onePL Image processing pipelines

```
template <formats Format>
using image_t = image<Format, std::uint8_t>;
```

```
// Size of source and destination buffers after color conversion
const sycl::range<2> src_size{ src_image_data.get_range() };
const auto          dst_size{ 2 * src_size };
const roi_rect      dst_roi_rect{ dst_size / 6, dst_size / 3 };
```

```
// Create queue
sycl::queue queue;
```

```
// Create allocator
shared usm_allocator t allocator{ queue };
```

```
// Source RGBA image data and destination images
```

```
image_t<formats::rgba> src_image{ src_image_data.get_pointer(), src_size };
```

```
image_t<formats::plane> gray_image{ src_size, allocator };
```

```
image_t<formats::plane> gray_image_roi{ gray_image, { src_size / 2 } };
```

```
image_t<formats::plane> sobel_image{ src_size, allocator };
```

```
image_t<formats::plane> sobel_image_roi{ sobel_image, { src_size * 2 / 3 } };
```

```
image_t<formats::plane> resized_image{ dst_size, allocator };
```

```
auto resized_image_roi = resized_image.get_roi(dst_roi_rect);
```


```
// Run pipeline: convert to grayscale -> sobel -> resize_lanczos -> resize_lanczos
```

```
(void)convert(queue, src_image, gray_image_roi);
```

```
(void)sobel_3x3(queue, gray_image, sobel_image);
```

```
(void)resize_lanczos(queue, sobel_image_roi, resized_image);
```

```
(void)resize_lanczos(queue, sobel_image_roi, resized_image_roi);
```



```
// Gray image data
```


```
// Gray image ROI
```

```
// Sobel image data
```

```
// Sobel image ROI
```

```
// Resized image data
```

```
// Resized image ROI
```



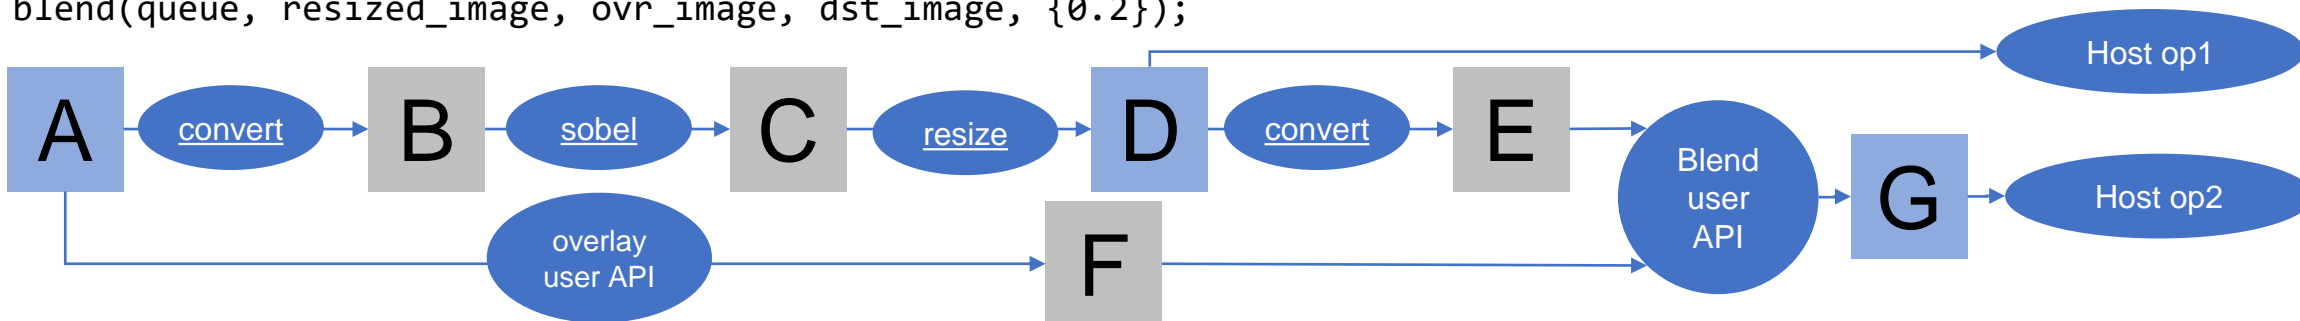
oneAPI

onePL Image processing pipelines


```
// Creation of I/O for pipeline steps
// Source RGBA image data and destination images (USM-based)
ipl::image<formats::rgba, uint8_t> src_image{ src_image_data.get_pointer(), src_size, image_alloc }; // A (image)
ipl::image<formats::plane, uint8_t> gray_image{ src_size, device_alloc }; // B (device USM)
ipl::image<formats::plane, uint8_t> sobel_image{ src_size, device_alloc }; // C (device USM)
ipl::image<formats::plane, uint8_t> resized_image{ dst_size, shared_alloc }; // D (shared USM)
ipl::image<formats::rgba, uint8_t> res_rgba_image{ dst_size, image_alloc }; // E (image)
ipl::image<formats::rgba, uint8_t> ovr_image{ dst_size, image_alloc }; // F (image)
ipl::image<formats::rgba, uint8_t> dst_image{ dst_size, image_alloc }; // G (image)
```

```
// Run pipeline: convert to grayscale -> sobel -> resize_lanczos
generate_overlay(queue, src_image, ovr_image.get_usm_pointer(), dst_size);
convert(queue, src_image, gray_image);
sobel_3x3(queue, gray_image, sobel_image);
resize_lanczos(queue, sobel_image, resized_image);
convert(queue, resized_image, res_rgba_image);
blend(queue, resized_image, ovr_image, dst_image, {0.2});
```

Legend	A	Input/output
	B	Temporary memory
	operation	Operation



onePL Image Abstraction




oneAPI


oneapi::ipl::image class is basic data abstraction for image data. onePL provides single abstraction over different memory types: host, device, shared and special GPU memory – textures.

```
template <formats Format, typename DataT, typename AllocatorT>
class image final
{
public:
    using image_impl_t::image_impl_t;
    using allocator_t = AllocatorT;
    using pixel_t = pixel_layout_t<DataT, Format>;

    image(const image_impl_t& image_impl);
    image(image_impl_t&& image_impl);
    auto operator=(const image_impl_t& image_impl)->image&;
    auto operator=(image_impl_t&& image_impl)->image&;

    auto get_whole_image() const->image;
    auto get_roi(const roi_rect& roi_rect) const->image;
};
```






onePL Memory Model

- onePL supports different type of memory – device, host, shared and special GPU texture (image) memory. Memory type is accessible explicitly via user allocation or via allocator passed to [`ipl::image`](#) constructor, if no external memory is provided. Allocator argument is a special tag. By default, allocator is selected as shared or texture depending on supported formats depending on device.

```
template <formats Format, typename DataT, typename AllocatorT =  
select_image_allocator_t<Format, DataT>>  
class image;
```


onePL Memory Model

- onePL supports 4 allocators which targets image memory allocation on host, device, shared or as a texture (image):

```
host_usm_allocator_t    host_allocator{ queue };    // memory is on host
shared_usm_allocator_t  shared_allocator{ queue }; // memory is on device and host
device_usm_allocator_t  device_allocator{ queue }; // memory is temporary (device only)
image_allocator_t       device_allocator{ queue }; // memory is texture (image)

ipl::image<formats::rgba,  uint8_t, host_usm_allocator_t> host_image{ size, host_allocator };
ipl::image<formats::rgba,  uint8_t, shared_usm_allocator_t> shared_image{ size, shared_allocator };
ipl::image<formats::rgba,  uint8_t, device_usm_allocator_t> device_image{ size, device_allocator };
ipl::image<formats::rgba,  uint8_t, image_allocator_t> texture_image{ size, image_allocator };

// doesn't work, since texture in SYCL 2020 supports only 4-channel images
ipl::image<formats::plane, uint8_t, image_allocator_t> texture_image{ size, image_allocator };
```




onePL Memory Model

- Image can be mapped over memory, but some cases shall be supported by device defined by [`sycl::queue`](#) object:

```
auto* device_ptr = sycl::malloc_device<std::uint8_t>(size[0] * size[1], queue);  
auto* shared_ptr = sycl::malloc_shared<std::uint8_t>(size[0] * size[1], queue);
```

```
ipl::image<formats::rgba, std::uint8_t, shared_usm_allocator_t> shared_image{ queue,  
device_ptr, size, shared_allocator };  
ipl::image<formats::rgba, std::uint8_t, device_usm_allocator_t> device_image{ queue,  
shared_ptr, size, device_allocator };
```

```
// if texture from device memory is supported by device, copy occurred, if not - would not work  
ipl::image<formats::plane, uint8_t, image_allocator_t> texture_image{ queue, device_ptr,  
size, image_allocator };
```



oneAPI

Next Steps


- All materials and minutes of meetings will be published on [GitHub](#) and will be available for the offline review (the offline feedback of invited TAB members will be also processed and discussed on next TAB meeting)
- The next technical discussion:
February 3rd (ww6)

Find more on <https://spec.oneapi.io/versions/latest/introduction.html#contribution-guidelines>
<https://github.com/oneapi-src/oneAPI-tab>

oneIPL Technical Advisory Board meetings


The goal is to provide the feedback and define future development of the specification.

First topics planned to discuss are at the table below, but it might be adjusted later.



oneAPI

Topic	Plan	Date
1) oneIPL overview	<ul style="list-style-type: none">1) Programming model2) Execution model3) Image processing pipelines4) Image data abstraction5) Memory model	
2) oneIPL Image data abstraction	<ul style="list-style-type: none">1) HW images and data formats and types coverage2) IPL image data abstraction3) Interoperability with USM4) Memory allocation and temporary images	February 3 rd , 2022
3) oneIPL Library design details	<ul style="list-style-type: none">1) Domains2) Reference code and optimized backends3) Error handling mechanism4) Interoperability with other oneAPI libraries	
4) oneIPL Functions overview	<ul style="list-style-type: none">1) ML oriented APIs for image preprocessing2) Data type support in the functions3) Color formats and conversions	



oneAPI

Resources

- <https://www.oneapi.io/spec/> - oneAPI Specification
- <https://spec.oneapi.io/oneapi/latest/index.html> - oneAPI specification (current version: v0.5)
- <https://github.com/oneapi-src/oneAPI-tab> - GitHub with oneAPI TAB materials
- <https://spec.oneapi.io/versions/latest/introduction.html#contribution-guidelines> - oneAPI Specification contribution guidelines