

Joint Matrix: A Unified SYCL Extension for Matrix Hardware Programming

Dounia Khaldi

Intel Corp.

oneAPI Language SIG Forum, June 07th, 2023

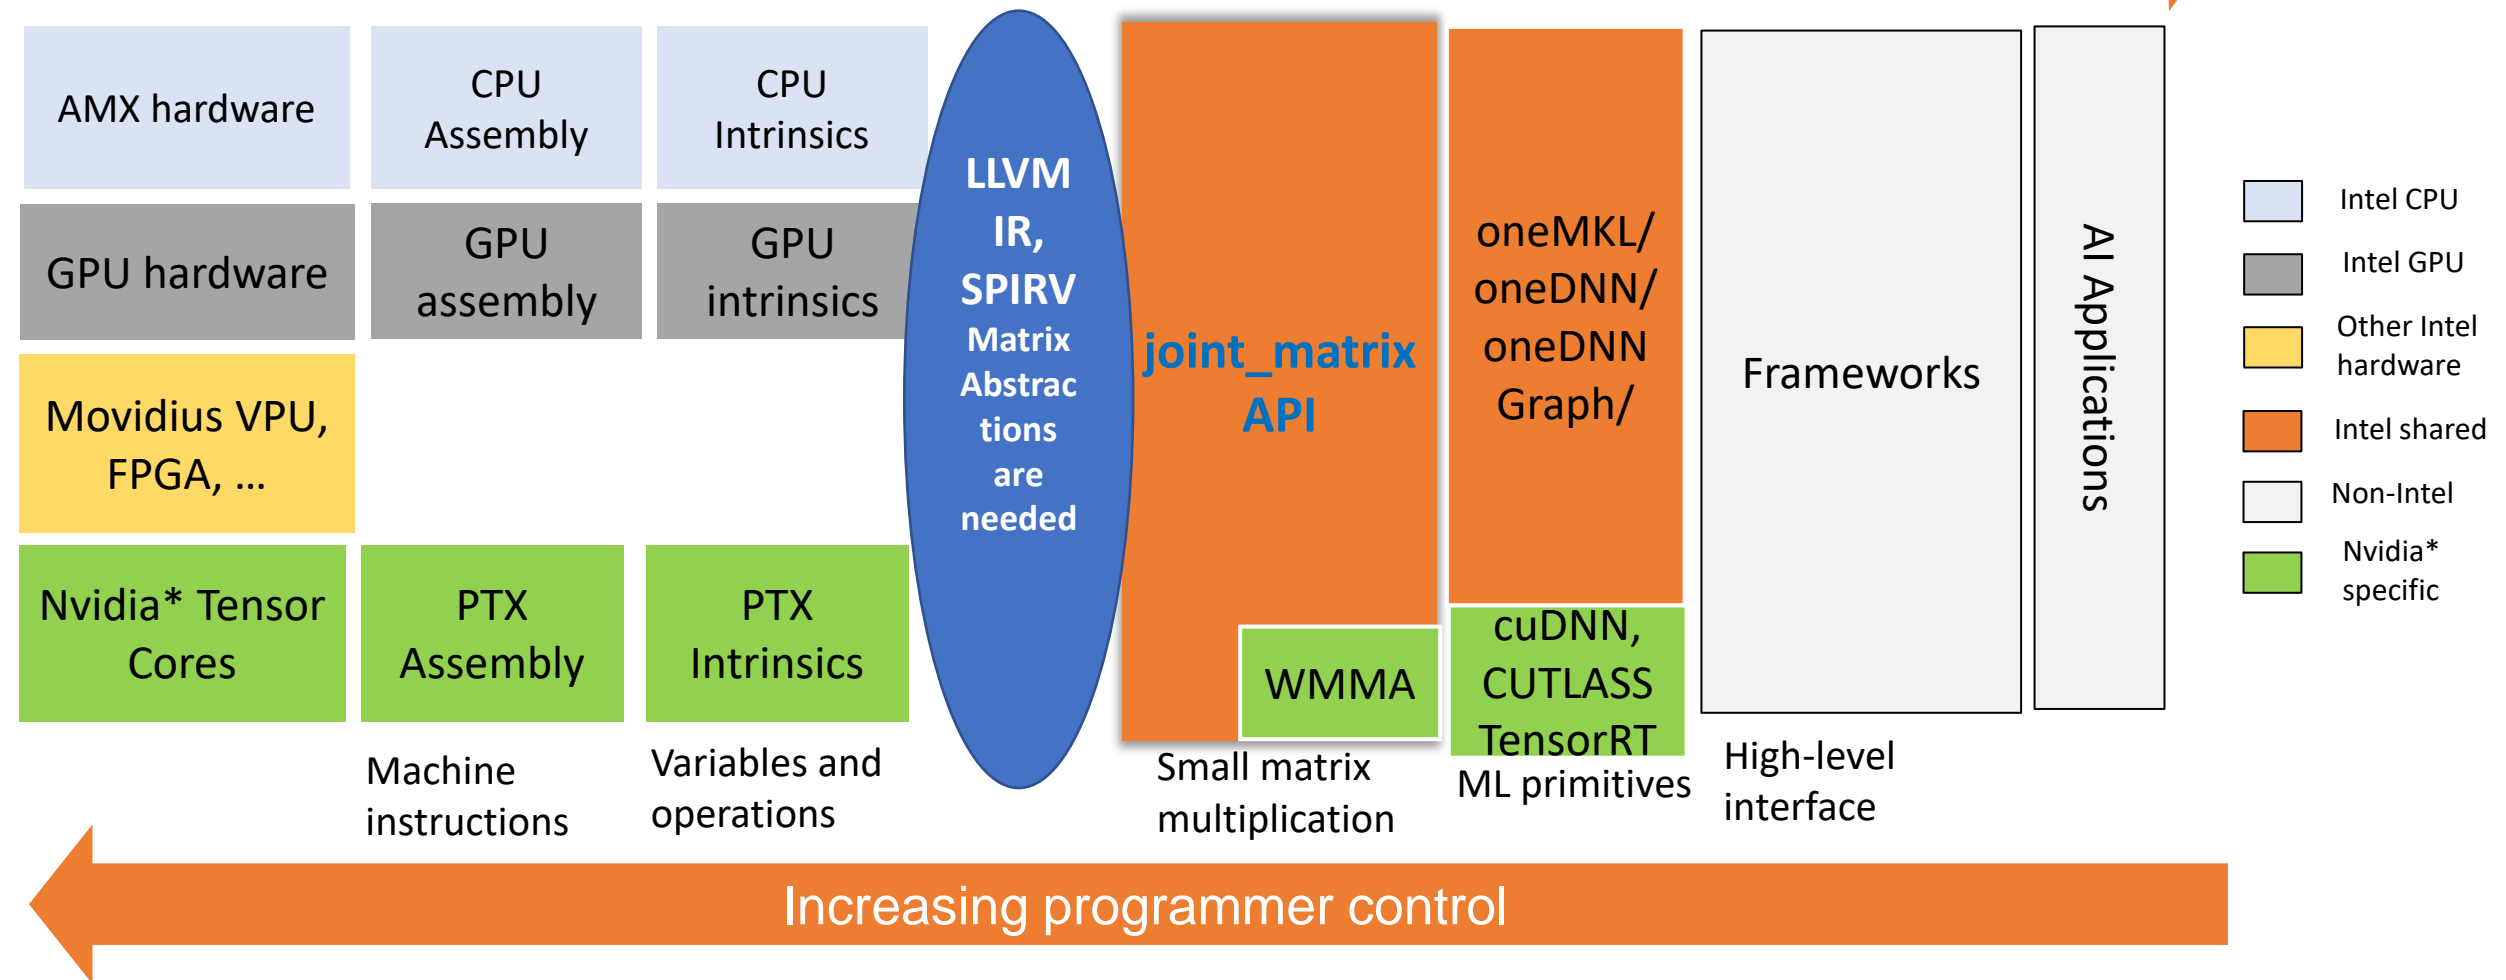
Collaborating with many colleagues from Intel and Codeplay

Executive Summary

- Goal
 - Deliver unified SYCL matrix interface across matrix hardware: Intel AMX, Intel XMV and Nvidia Tensor Cores
 - Programmer productivity: Allow the customer to express their applications for matrix hardware with minimal changes
 - Performance: Maps directly to low-level intrinsics/assembly for maximum performance
- Status
 - Implementation: Unified interface is part of oneAPI 2023.1 release
 - Current Users: Code porting from CUDA wmma, MLIR SPIRV-based joint matrix code generation

Programming Abstractions for Matrix Computing

Increasing level of Abstraction



Lead Users

Performance portability across all hardware without extra effort of optimizations for specific hardware



AI Scientists

- New operations such as tensor contractions, BRGEMM, quantized gemm, fused operations



Library Developers


- Different DNN and BLAS libraries



Matrix Hardware

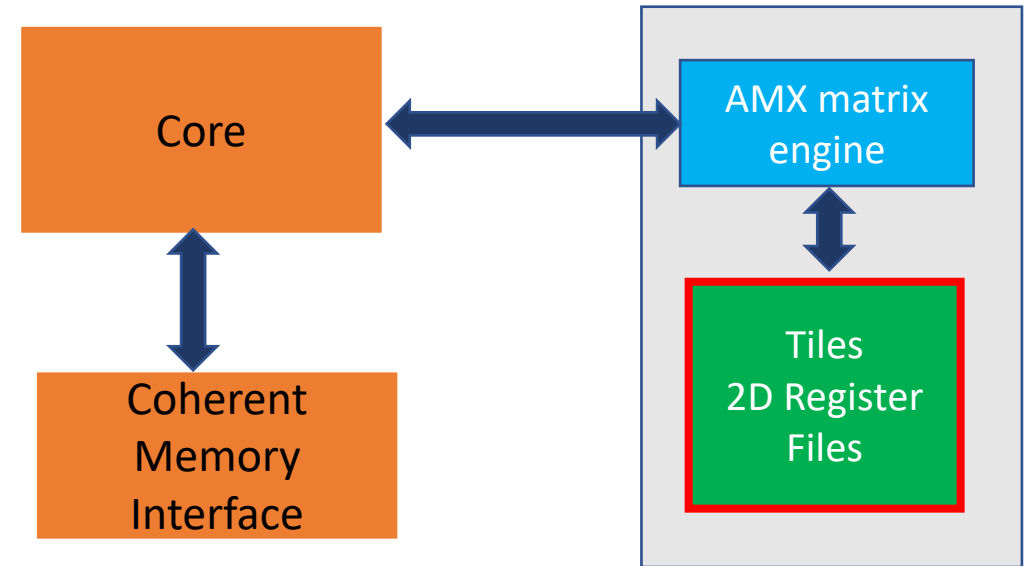
Intel XMX in Intel® Data Center GPU Max Series

- Code-named Ponte Vecchio (PVC)
- Xe-HPC 2-Stack Ponte Vecchio GPU
- Each Xe-Stack has 4 slices
- Xe-slice contains 16 Xe-core
- An Xe-core contains 8 vector and 8 matrix engines



Intel AMX High-Level Architecture

- Intel® Xeon® processor codenamed Sapphire Rapids
- Intel AMX, an Intel x86 extension for multiplication of matrices of bf16/int8 elements



SYCL Joint Matrix Extension

SYCL Matrix Extension

Namespace

namespace sycl::ext::oneapi::experimental::matrix

- New matrix data type with group scope
- Defined with a specified type, use (a, b, accumulator), size, and layout

```
template <typename Group, typename T, use Use, size_t  
    Rows, size_t Cols, layout Layout = layout::dynamic>  
struct joint_matrix;  
enum class use { a, b, accumulator};
```

- Separate memory operations from the compute
- enum class layout {row_major, col_major, dynamic};
- Group execution scope → *joint, Group as argument*

- **joint_matrix_fill**(Group g, joint_matrix<>&dst, T v);
- void **joint_matrix_load**(Group g, joint_matrix<>dst, T *base, unsigned stride, Layout layout);
- void **joint_matrix_store**(Group g, joint_matrix<>src, T *base, unsigned stride, Layout layout);

- Multiply and add
- Element-wise ops
- Extensible to add more operations

- joint_matrix<> **joint_matrix_mad**(Group g, joint_matrix<>A, joint_matrix<>B, joint_matrix<>C);
- void **joint_matrix_apply**(Group g, joint_matrix<>A, F&& func);

SYCL

joint_matrix

Example

```
using namespace sycl::ext::oneapi::experimental::matrix;
queue q;
range<2> G = {M/tM, N/tN * SG_SIZE};
range<2> L = {1, SG_SIZE};
auto bufA = sycl::buffer{memA, sycl::range{M*K}};
auto bufB = sycl::buffer{memB, sycl::range{K*N}};
auto bufC = sycl::buffer{memC, sycl::range{M*N}};
q.submit([&](sycl::handler& cgh) {
    auto accA = sycl::accessor{bufA, cgh, sycl::read_only};
    auto accB = sycl::accessor{bufB, cgh, sycl::read_only};
    auto accC = sycl::accessor{bufC, cgh, sycl::read_write};
    cgh.parallel_for(nd_range<2>(G, L), [=](nd_item<2> item) {
        const auto sg_startx = item.get_global_id(0) - item.get_local_id(0);
        const auto sg_starty = item.get_global_id(1) - item.get_local_id(1);
        sub_group sg = item.get_sub_group();
        joint_matrix<sub_group, int8_t, use::a, tM, tK, layout::row_major> subA;
        joint_matrix<sub_group, int8_t, use::b, tK, tN, layout::row_major> subB;
        joint_matrix<sub_group, int32_t, use::accumulator, tM, tN> subC;
        joint_matrix_fill(sg, subC, 0);
        for (int k = 0; k < K; k += tk) {
            joint_matrix_load(sg, subA, accA.get_pointer() + sg_startx * tM * K + k, K);
            joint_matrix_load(sg, subB, accB.get_pointer() + k * N + sg_starty, N);
            subC = joint_matrix_mad(sg, subA, subB, subC);
        }
        joint_matrix_apply(sg, subC, [=](T &x) { Relu(x); });
        joint_matrix_store(sg, subC.get_pointer(), accC + sg_startx * tM * N + sg_starty,
N, row_major);
    });
});
q.wait;
```

oneAPI 2023.1: One Joint Matrix Code to Run on Intel AMX, Intel XMX and Nvidia* Tensor Cores

```
joint_matrix<sub_group, int8_t, use::a, tM, tK, layout::row_major> subA;  
joint_matrix<sub_group, int8_t, use::b, tK, tN, layout::row_major> subB;  
joint_matrix<sub_group, int32_t, use::accumulator, tM, tN> subC;  
sub_group sg = item.get_sub_group();  
joint_matrix_fill(sg, subC, 0);  
for (int k = 0; k < K; k += tK) {  
    joint_matrix_load(sg, subA, accA.get_pointer()+ sg_startx * tM * K + k, K);  
    joint_matrix_load(sg, subB, accB.get_pointer()+ k * N + sg_starty/SG_SIZE*tN, N);  
    subC = joint_matrix_mad(sg, subA, subB, subC);  
}  
joint_matrix_apply(sg, subC, [=](T x) { x *= alpha; });  
joint_matrix_store(sg, subC, accC.get_pointer() + sg_startx * tM * N + sg_starty/SG_SIZE*tN, N, layout::row_major);
```

Intel
CPUs

Intel
GPUs

Nvidia*
GPUs

SYCL Matrix Extension: Intel Specific Features

SYCL joint_matrix Indexing with Coordinates

- Element wise ops that apply to a set of elements of the matrix → Mapping is required
- Example: Quantization Calculations
- $A*B + \text{sum_rows_A} + \text{sum_cols_B} + \text{scalar_zero_point}$
- `sum_rows_A` returns a single row of A

```
using namespace sycl::ext::intel::experimental::matrix;

void sum_rows_A(joint_matrix<T, rows, cols>& subA)
{
    joint_matrix_apply(sg, subA, [=](T &val, size_t row, size_t col) {
        global_row = row + global_idx * rows;
        sum_local_rows[global_row] += val;
    });
}
```

Matrix Query Interface

AMX Supported Combinations

A type	Btype	Ctype	M	N	K
(u)int8_t	(u)int8_t	int32_t	≤ 16	≤ 16	≤ 64
bf16	bf16	float	≤ 16	≤ 16	≤ 32

Intel XMX Supported Combinations

A type	Btype	Ctype	M	N	K
(u)int8_t	(u)int8_t	int32_t	<=8	8 (ATS-M) 16 (PVC)	32
fp16	fp16	float	<=8	8 (ATS-M) 16 (PVC)	16
bf16	bf16	float	<=8	8 (ATS-M) 16 (PVC)	16
tf32	tf32	float	<=8	16 (PVC)	

Nvidia* Tensor Cores Supported Combinations

A type	Btype	Accumulator type	M	N	K
half	half	float	16	16	16
			32	8	16
			8	32	16
half	half	half	16	16	16
			32	8	16
			8	32	16
bfloat16	bfloat16	float	16	16	16
			32	8	16
			8	32	16
tf32	tf32	float	16	16	8
(u)int8_t	(u)int8_t	int32_t	16	16	16
			32	8	16
			8	32	16

Matrix Query

Static Query

Provide a default shape if user does not provide a combination in a *constexpr* way

```
namespace sycl::ext::oneapi::experimental::matrix
template<sycl::ext::oneapi::experimental::architecture Dev, typename Ta, typename Tb,
        typename Taccumulator>
struct matrix_params {
    static constexpr size_t M = /* implementation defined */;
    static constexpr size_t N = /* implementation defined */;
    static constexpr size_t K = /* implementation defined */;
    template <typename Group, layout Layout>
    using joint_matrix_a = joint_matrix<Group, Ta, use::a, M, K, Layout>;
    template <typename Group, layout Layout>
    using joint_matrix_b = joint_matrix<Group, Tb, use::b, K, N, Layout>;
    template <typename Group>
    using joint_matrix_accumulator = joint_matrix<Group, Taccumulator, use::accumulator, M, N>;
};
```

Runtime Query

Tell the set of supported matrix sizes and types on this device
(Not implemented yet)

```
namespace sycl::ext::oneapi::experimental::info::device::matrix
std::vector<combination> combinations =
    device.get_info<info::device::matrix::combinations>();
for (int i = 0; sizeof(combinations); i++) {
    if (Ta == combinations[i].atype && Tb == combinations[i].btype &&
        Tc == combinations[i].ctype && Td == combinations[i].dtype) {
        // joint matrix GEMM kernel can be called using these sizes
        joint_matrix_gemm(combinations[i].msize, combinations[i].nsize, combinations[i].ksize);
    }
}
```

SYCL *joint_matrix* Using the Default Query

```
using namespace sycl::ext::oneapi::experimental::matrix;
using myparams = matrix_params<sycl::ext::oneapi::experimental::architecture::intel_gpu_pvc,
                                int8_t, int8_t, int>;

constexpr int tM = myparams::M;
constexpr int tN = myparams::N;
constexpr int tK = myparams::K;
range<2> G = {M/tM, N/tN * SG_SIZE};
range<2> L = {1, SG_SIZE};
// buffers
q.submit([&](sycl::handler& cgh) {
    // accessors
    cgh.parallel_for(nd_range<2>(G, L), [=](nd_item<2> item) {
        const auto sg_startx = item.get_global_id(0) - item.get_local_id(0);
        const auto sg_starty = item.get_global_id(1) - item.get_local_id(1);
        sub_group sg = item.get_sub_group();
        myparams::joint_matrix_a<sub_group, layout::row_major> tA;
        myparams::joint_matrix_b<sub_group, layout::row_major> tB;
        myparams::joint_matrix_accumulator<sub_group> tC;
        joint_matrix_fill(sg, tC, 0);
        for (int k = 0; k < K; k += tk) {
            joint_matrix_load(sg, tA, memA + sg_startx * tM * K + k, K);
            joint_matrix_load(sg, tB, memB + k * N + sg_starty, N);
            tC = joint_matrix_mad(sg, tA, tB, tC);
        }
        joint_matrix_store(sg, tC, memC + sg_startx * tM * N + sg_starty, N, row_major);
    });
});
```

SYCL joint_matrix

```
// inputA is MxK, inputB is KxN, inputC is MxN
#define tM=16 tN=16 tK=16

void gemm(size_t global_idx, size_t global_idy, size_t local_idx, size_t local_idy, sub_group sg) {

    joint_matrix<sub_group, half, use::a, tM, tK, row_major> matA;
    joint_matrix<sub_group, half, use::b, tK, tN, row_major> matB;
    joint_matrix<sub_group, float, use::accumulator, tM, tN> matC;

    const auto sg_startx = global_idx - local_idx;
    const auto sg_starty = global_idy - local_idy;

    joint_matrix_fill(matC, 0.0f);

    for (int step = 0; step < K; step += tK) {

        uint AStart = sg_startx * tM * K + step;
        uint BStart = step * N + sg_starty;
        joint_matrix_load(sg, matA, inputA + AStart, K);
        joint_matrix_load(sg, matB, inputB + BStart, N);
        matC = joint_matrix_mad(sg, matA, matB, matC);

    }

    joint_matrix_apply(sg, matC, [=](T& x) { x *= alpha; });

    joint_matrix_store(sg, matC, output + sg_startx * tM * N + sg_starty, N, row_major);

}
```



CUDA Fragments

```
// inputA is MxK, inputB is KxN, inputC is MxN
#define tM=16 tN=16 tK=16

__global__ void wmma_ker(blockidx) {

    fragment<matrix_a, 16, 16, 16, half, col_major> a_frag;
    fragment<matrix_b, 16, 16, 16, half, row_major> b_frag;
    fragment<accumulator, 16, 16, 16, float> c_frag;

    uint row = (blockIdx%x - 1)*tM + 1
    uint col = (blockIdx%y - 1)*tN + 1

    fill_fragment(c_frag, 0.0f);

    for (uint step = 0; step < K; step += matrixDepth) {

        uint AStart = row * rowStrideA + step;
        uint BStart = col * colStrideB + step;

        load_matrix_sync(matA, inputA + AStart, K);

        load_matrix_sync(matB, inputB + BStart, N);

        mma_sync(matC, matA, matB, matC);

    }

    for(int t=0; t<matC.num_elements; t++)

        matC.x[t] *= alpha;

    store_matrix_sync(inputC+row*N+col, matC, N, mem_row_major);

}
```

Current Users of Joint Matrix

Code migration from wmma samples

- https://github.com/wzsh/wmma_tensorcore_sample/tree/master/matrix_wmma/matrix_wmma
- https://github.com/NVIDIA/cuda-samples/tree/master/Samples/3_CUDA_Features/cudaTensorCoreGemm

Porting code from wmma to joint_matrix

- Porting an earthquake simulation code that makes direct use of the tensor cores through wmma from CUDA and wmma to SYCL and joint matrix

SYCL-DNN – By CodePlay

- Using joint_matrix for enabling Nvidia Tensor Cores in SYCL-DNN

SYCL-BLAS – By CodePlay

- Using joint_matrix for enabling Nvidia Tensor Cores in SYCL-BLAS GEMM

SPIRV MLIR Dialect

- XMX Support using MLIR SPIRV dialect by adding SPIRV joint_matrix

Conclusion and Next Steps

- Full support of SYCL joint matrix extension on Intel AMX, Intel XMX, and Nvidia Tensor Cores
- Extensions to LLVM IR and SPIRV
- Effective usage in MLIR integration and CUDA code migration
- Next steps:
 - Standardization of SYCL joint matrix to Khronos SYCL
 - Standardization of SPIRV joint matrix to Khronos SPIRV
- Contributions/feedback are welcome

Legal Notices & Disclaimers

Intel technologies may require enabled hardware, software or service activation.

No product or component can be absolutely secure.

Your costs and results may vary.

© Intel Corporation. Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries. Other names and brands may be claimed as the property of others.